

# Inlined Code Generation for Smalltalk

Dave Mason and Daniel Franklin

Toronto Metropolitan University

©2024 Dave Mason



Toronto  
Metropolitan  
University

# Why inline?

- obviate need for sends/calls
- access variables more efficiently
- use registers more effectively - modern hardware
- larger body of code for optimizers
- hierarchy of need:
  - static procedural
  - static vtables
  - dynamic dispatch - Java interfaces, dynamic languages

# Why inline Smalltalk?

- dynamic dispatch, and also:
  - very small methods are encouraged
  - inheritance is heavily used
  - control structures are message sends
  - non-local returns are common

# Smalltalk Method Dispatch

- 1 set T to the class of receiver, set S to selector
  - 2 look up S in T methods - if found execute the method
  - 3 else if T has a superclass, set T to the superclass and continue from step 2
  - 4 else set T to the class of receiver, set S to `doesNotUnderstand:` with the original message as the parameter, and continue from step 2
- same as Java except DNU and offsets are statically detected (except interfaces)

## What's known?

- Deutsch-Schiffman dispatch has been used since the 1980s
- about a dozen “special” control-flow selectors are inlined
- caches ameliorate the cost - Polymorphic Inline Cache
- Self in the early 1990s (and subsequently Javascript) compiled JITed code using per-target dispatch tables and inlining
- inlining almost exclusively in native-code JITs

# What's new?

- Zag is new Smalltalk implementation
- two interchangeable execution modes - threaded (for debugging), and native (for speed)
- compiles methods to concrete classes (like Self & Javascript) - so `self` and `super` are usually exact
- compilation is “on demand” rather than JIT or AOT (only receiver class considered) - tiered:
  - threaded, no inlining
  - threaded, inlining
  - native version of previous

## ... Inlining

- no *ad hoc* selector inlining
- inlining is semantic - independent of target code - debugging is of the same code
- inlining is done breadth-first
- can be more intentional about inlining
- can be filtered to manage code explosion, possible strategies:
  - n-levels
  - maximum number of inlinings
  - inline until all possible block closures have been inlined
  - inline only sends with block closure parameters
  - incrementally inline - additional rounds

# Inlining Opportunities

- Send to `self`, `super` or known type
- Send to `self`, `super` or known type when can't inline
- Recursive send to `self`, `super` or known type
- Recursive tail-call send to `self`, `super` or known type
- Send to `self`, `super` or known type where the method is primitive
- Send `value`, `value:`, etc. to a literal `BlockClosure` - *safe*
- Send where there are few implementations of a method - *not safe*
- Send where the target is the result of a comparison primitive - *safe*



# Conclusions

- there are advantages to semantic inlining
- we will report back when we have some actual benchmark results
- <https://github.com/Zag-Research/Zag-Smalltalk>

# Questions?

we are hiring and have (many) open faculty positions in Toronto

`dmason@torontomu.ca`